

Dragon_Real_Estates

January 5, 2021

0.1 Dragon Real Estates - Price Predictor

```
[1]: import pandas as pd
```

```
[2]: housing = pd.read_csv('housing.csv')
```

```
[3]: housing.head()
```

```
[3]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
1	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
2	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
3	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	
4	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	9.14	21.6
1	392.83	4.03	34.7
2	394.63	2.94	33.4
3	396.90	5.33	36.2
4	394.12	5.21	28.7

```
[4]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 505 entries, 0 to 504
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        505 non-null   float64
1   ZN          505 non-null   float64
2   INDUS       505 non-null   float64
3   CHAS        505 non-null   int64
4   NOX         505 non-null   float64
5   RM          500 non-null   float64
6   AGE         505 non-null   float64
7   DIS         505 non-null   float64
8   RAD         505 non-null   int64
9   TAX         505 non-null   int64
```

```

10 PTRATIO  505 non-null    float64
11 B        505 non-null    float64
12 LSTAT    505 non-null    float64
13 MEDV     505 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.4 KB

```

```
[5]: housing['CHAS'].value_counts()
```

```

[5]: 0    470
     1     35
     Name: CHAS, dtype: int64

```

```
[6]: housing.describe()
```

```

[6]:
      count    CRIM      ZN      INDUS      CHAS      NOX      RM  \
count  505.000000  505.000000  505.000000  505.000000  505.000000  500.000000
mean    3.620667  11.350495  11.154257    0.069307    0.554728    6.285460
std     8.608572  23.343704   6.855868    0.254227    0.115990    0.706416
min     0.009060   0.000000   0.460000    0.000000    0.385000    3.561000
25%     0.082210   0.000000   5.190000    0.000000    0.449000    5.883000
50%     0.259150   0.000000   9.690000    0.000000    0.538000    6.208500
75%     3.678220  12.500000  18.100000    0.000000    0.624000    6.629250
max    88.976200 100.000000  27.740000    1.000000    0.871000    8.780000

      count    AGE      DIS      RAD      TAX      PTRATIO      B  \
count  505.000000  505.000000  505.000000  505.000000  505.000000  505.000000
mean    68.581584   3.794459   9.566337  408.459406   18.461782  356.594376
std    28.176371   2.107757   8.707553  168.629992   2.162520   91.367787
min     2.900000   1.129600   1.000000  187.000000  12.600000   0.320000
25%    45.000000   2.100000   4.000000  279.000000  17.400000  375.330000
50%    77.700000   3.199200   5.000000  330.000000  19.100000  391.430000
75%    94.100000   5.211900  24.000000  666.000000  20.200000  396.210000
max   100.000000  12.126500  24.000000  711.000000  22.000000  396.900000

      count    LSTAT      MEDV
count  505.000000  505.000000
mean    12.668257   22.529901
std     7.139950    9.205991
min     1.730000    5.000000
25%     7.010000   17.000000
50%    11.380000   21.200000
75%    16.960000   25.000000
max    37.970000   50.000000


```

```
[7]: %matplotlib inline
```

```
[8]: import matplotlib.pyplot as plt
```

```
[9]: housing.hist(bins=50, figsize=(20,15))
```

```
[9]: array([[<AxesSubplot:title={'center':'CRIM'}>,
  <AxesSubplot:title={'center':'ZN'}>,
  <AxesSubplot:title={'center':'INDUS'}>,
  <AxesSubplot:title={'center':'CHAS'}>],
[<AxesSubplot:title={'center':'NOX'}>,
  <AxesSubplot:title={'center':'RM'}>,
  <AxesSubplot:title={'center':'AGE'}>,
  <AxesSubplot:title={'center':'DIS'}>],
[<AxesSubplot:title={'center':'RAD'}>,
  <AxesSubplot:title={'center':'TAX'}>,
  <AxesSubplot:title={'center':'PTRATIO'}>,
  <AxesSubplot:title={'center':'B'}>],
[<AxesSubplot:title={'center':'LSTAT'}>,
  <AxesSubplot:title={'center':'MEDV'}>, <AxesSubplot:>,
  <AxesSubplot:>]], dtype=object)
```



0.2 Train Test Splitting

```
[10]: import numpy as np
      # for learning purupose
      def split_train_test(data, test_ratio):
          np.random.seed(42)
          shuffled = np.random.permutation(len(data))
          test_set_size = int(len(data) * test_ratio)
          test_indices = shuffled[:test_set_size]
          train_indices = shuffled[test_set_size:]
          return data.iloc[train_indices], data.iloc[test_indices]
```

```
[11]: train_set, test_set = split_train_test(housing, 0.2)
```

```
[12]: # print(f"Rows in train set: {len(train_set)}\nRows in test set:
      ↪{len(test_set)}\n")
```

```
[13]: from sklearn.model_selection import train_test_split
      train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)
      print(f"Rows in train set: {len(train_set)}\nRows in test set:
      ↪{len(test_set)}\n")
```

Rows in train set: 404

Rows in test set: 101

```
[14]: from sklearn.model_selection import StratifiedShuffleSplit
      split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
      for train_index, test_index in split.split(housing, housing['CHAS']):
          strat_train_set = housing.loc[train_index]
          strat_test_set = housing.loc[test_index]
```

```
[15]: strat_train_set['CHAS'].value_counts()
```

```
[15]: 0    376
      1    28
      Name: CHAS, dtype: int64
```

```
[16]: # 376/28
```

```
[17]: # strat_test_set['CHAS'].value_counts()
```

```
[18]: # 94/7
```

```
[19]: housing = strat_train_set.copy()
```

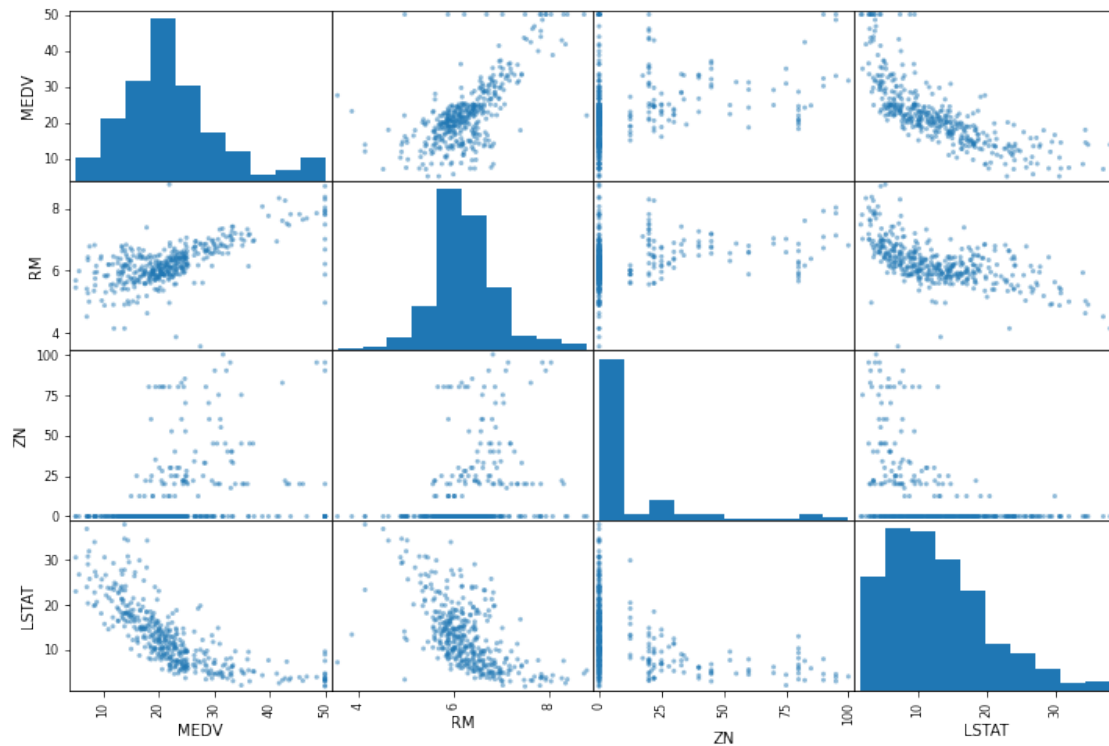
0.3 Looking for Correlations

```
[20]: corr_matrix = housing.corr()
      corr_matrix['MEDV'].sort_values(ascending=False)
```

```
[20]: MEDV      1.000000
      RM       0.661599
      B       0.344609
      ZN       0.329206
      DIS      0.231680
      CHAS     0.215042
      RAD     -0.362619
      AGE     -0.378913
      CRIM    -0.397993
      NOX     -0.421815
      TAX     -0.441617
      INDUS   -0.448303
      PTRATIO -0.486045
      LSTAT   -0.739129
      Name: MEDV, dtype: float64
```


```
[21]: from pandas.plotting import scatter_matrix
      attributes = ['MEDV', 'RM', 'ZN', 'LSTAT']
      scatter_matrix(housing[attributes], figsize = (12,8))
```

```
[21]: array([[<AxesSubplot:xlabel='MEDV', ylabel='MEDV'>,
              <AxesSubplot:xlabel='RM', ylabel='MEDV'>,
              <AxesSubplot:xlabel='ZN', ylabel='MEDV'>,
              <AxesSubplot:xlabel='LSTAT', ylabel='MEDV'>],
             [<AxesSubplot:xlabel='MEDV', ylabel='RM'>,
              <AxesSubplot:xlabel='RM', ylabel='RM'>,
              <AxesSubplot:xlabel='ZN', ylabel='RM'>,
              <AxesSubplot:xlabel='LSTAT', ylabel='RM'>],
             [<AxesSubplot:xlabel='MEDV', ylabel='ZN'>,
              <AxesSubplot:xlabel='RM', ylabel='ZN'>,
              <AxesSubplot:xlabel='ZN', ylabel='ZN'>,
              <AxesSubplot:xlabel='LSTAT', ylabel='ZN'>],
             [<AxesSubplot:xlabel='MEDV', ylabel='LSTAT'>,
              <AxesSubplot:xlabel='RM', ylabel='LSTAT'>,
              <AxesSubplot:xlabel='ZN', ylabel='LSTAT'>,
              <AxesSubplot:xlabel='LSTAT', ylabel='LSTAT'>]], dtype=object)
```



```
[22]: housing.plot(kind='scatter', x='RM', y='MEDV', alpha=0.8)
```

```
[22]: <AxesSubplot:xlabel='RM', ylabel='MEDV'>
```



1 Trying out Attribute Combinations

```
[23]: housing['TAXRM'] = housing['TAX']/housing['RM']
```

```
[24]: housing.head()
```

```
[24]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
254	0.03548	80.0	3.64	0	0.392	5.876	19.1	9.2203	1	315	
348	0.02899	40.0	1.25	0	0.429	6.939	34.5	8.7921	1	335	
476	15.02340	0.0	18.10	0	0.614	5.304	97.3	2.1007	24	666	
321	0.35114	0.0	7.38	0	0.493	6.041	49.9	4.7211	5	287	
326	0.24103	0.0	7.38	0	0.493	6.083	43.7	5.4159	5	287	


	PTRATIO	B	LSTAT	MEDV	TAXRM
254	16.4	395.18	9.25	20.9	53.607897
348	19.7	389.85	5.89	26.6	48.277850
476	20.2	349.48	24.91	12.0	125.565611
321	19.6	396.90	7.70	20.4	47.508691
326	19.6	396.90	12.79	22.2	47.180667

```
[25]: corr_matrix = housing.corr()
      corr_matrix['MEDV'].sort_values(ascending=False)
```

```
[25]: MEDV      1.000000
      RM       0.661599
      B       0.344609
      ZN       0.329206
      DIS      0.231680
      CHAS     0.215042
      RAD     -0.362619
      AGE     -0.378913
      CRIM    -0.397993
      NOX     -0.421815
      TAX     -0.441617
      INDUS   -0.448303
      PTRATIO -0.486045
      TAXRM   -0.509117
      LSTAT   -0.739129
      Name: MEDV, dtype: float64
```

```
[26]: housing.plot(kind='scatter', x='TAXRM', y='MEDV', alpha=0.8)
```

```
[26]: <AxesSubplot:xlabel='TAXRM', ylabel='MEDV'>
```



```
[27]: housing = strat_train_set.drop('MEDV', axis=1)
housing_labels = strat_train_set["MEDV"].copy()
```

1.1 Missing Attributes

```
[28]: # to take care of missing attributes
# 1. Get rid off the missing data points
# 2. Get rid off the whole attribute
# 3. Set the value to some value (0, mean or meadian)
```

```
[29]: a= housing.dropna(subset=['RM']) #option 1
a.shape
# note that the original housing df will remain unchanged
```

```
[29]: (401, 13)
```

```
[30]: housing.drop('RM', axis=1).shape #option 2
# note that there is no RM column and also note that the original housing df
↳ will remain unchanged
```

```
[30]: (404, 12)
```

```
[31]: median = housing['RM'].median() #compute median for option 3
```



```
[32]: housing['RM'].fillna(median) #option 3
      # note that the original housing df will remain unchanged
```

```
[32]: 254    5.876
      348    6.939
      476    5.304
      321    6.041
      326    6.083
      ...
      154    6.152
      423    5.565
      98     7.416
      455    5.976
      215    5.888
      Name: RM, Length: 404, dtype: float64
```

```
[33]: housing.shape
```

```
[33]: (404, 13)
```

```
[34]: housing.describe() # before we started filling missing attributes
```

```
[34]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	404.000000	404.000000	404.000000	404.000000	404.000000	401.000000
mean	3.680733	10.189356	11.305965	0.069307	0.557274	6.251918
std	8.249705	21.930822	6.817698	0.254290	0.116503	0.691261
min	0.009060	0.000000	0.740000	0.000000	0.385000	3.561000
25%	0.090060	0.000000	5.190000	0.000000	0.452000	5.874000
50%	0.290250	0.000000	9.900000	0.000000	0.538000	6.176000
75%	3.694070	3.125000	18.100000	0.000000	0.625750	6.606000
max	73.534100	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000
mean	68.548020	3.778549	9.702970	411.428218	18.502723	353.522649
std	28.433028	2.125958	8.754489	168.237476	2.117437	95.111003
min	2.900000	1.129600	1.000000	187.000000	13.000000	0.320000
25%	44.850000	2.070275	4.000000	284.000000	17.400000	374.237500
50%	77.500000	3.167500	5.000000	336.000000	19.050000	390.940000
75%	94.600000	5.104475	24.000000	666.000000	20.200000	396.157500
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT
count	404.000000
mean	12.833292
std	7.199418
min	1.730000

```

25%      7.362500
50%     11.570000
75%     16.977500
max      37.970000

```

```

[35]: from sklearn.impute import SimpleImputer
      imputer = SimpleImputer(strategy = 'median')
      imputer.fit(housing)

```

```

[35]: SimpleImputer(strategy='median')

```

```

[36]: imputer.statistics_

```

```

[36]: array([2.9025e-01, 0.0000e+00, 9.9000e+00, 0.0000e+00, 5.3800e-01,
          6.1760e+00, 7.7500e+01, 3.1675e+00, 5.0000e+00, 3.3600e+02,
          1.9050e+01, 3.9094e+02, 1.1570e+01])

```

```

[37]: x = imputer.transform(housing)

```

```

[38]: housing_tr = pd.DataFrame(x, columns= housing.columns)

```

```

[39]: housing_tr.describe() #now rm count will increase from 500 to 501

```

```

[39]:
      count  CRIM      ZN      INDUS      CHAS      NOX      RM  \
count  404.000000  404.000000  404.000000  404.000000  404.000000  404.000000
mean    3.680733   10.189356   11.305965    0.069307    0.557274    6.251354
std     8.249705   21.930822    6.817698    0.254290    0.116503    0.688714
min     0.009060    0.000000    0.740000    0.000000    0.385000    3.561000
25%     0.090060    0.000000    5.190000    0.000000    0.452000    5.874750
50%     0.290250    0.000000    9.900000    0.000000    0.538000    6.176000
75%     3.694070    3.125000   18.100000    0.000000    0.625750    6.604500
max    73.534100  100.000000   27.740000    1.000000    0.871000    8.780000

```

```

      count  AGE      DIS      RAD      TAX      PTRATIO      B  \
count  404.000000  404.000000  404.000000  404.000000  404.000000  404.000000
mean    68.548020    3.778549    9.702970  411.428218   18.502723  353.522649
std    28.433028    2.125958    8.754489  168.237476    2.117437   95.111003
min     2.900000    1.129600    1.000000  187.000000   13.000000    0.320000
25%    44.850000    2.070275    4.000000  284.000000   17.400000  374.237500
50%    77.500000    3.167500    5.000000  336.000000   19.050000  390.940000
75%    94.600000    5.104475   24.000000  666.000000   20.200000  396.157500
max   100.000000   12.126500   24.000000  711.000000   22.000000  396.900000

```

```

      count  LSTAT
count  404.000000
mean    12.833292
std     7.199418
min     1.730000

```

25%	7.362500
50%	11.570000
75%	16.977500
max	37.970000

1.2 ScikitLearn Design

Primarily, three types of objects 1. Estimators- It estimates some parameter based on a dataset. Eg imputer. It has a fit method and transform method.

1.3 Feature Scaling

Primarily, two types of feature scaling methods: 1. Min-Max scaling (Normalization) (value - min)/(max - min) ranges from 0 to 1 Sklearn provides a class MinMaxScaler for this

2. Standardization (value - mean)/std Sklearn provides a class called Standard Scaler for this

1.4 Creating a Pipeline

```
[40]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      my_pipeline = Pipeline([
          ('imputer', SimpleImputer(strategy = 'median')),
          # ..... add as many as you want in your pipeline
          ('std_scaler', StandardScaler()),
      ])
```

```
[41]: housing_num_tr = my_pipeline.fit_transform(housing)
```

```
[42]: housing_num_tr.shape
```

```
[42]: (404, 13)
```

1.5 Selecting a desired model for Dragon Real Estates

```
[44]: from sklearn.linear_model import LinearRegression
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.ensemble import RandomForestRegressor
      # model = LinearRegression()
      # model = DecisionTreeRegressor()
      model = RandomForestRegressor()
      model.fit(housing_num_tr, housing_labels)
```

```
[44]: RandomForestRegressor()
```

```
[45]: some_data = housing.iloc[:5]
```

```
[46]: some_labels = housing_labels.iloc[:5]
```

```
[47]: prepared_data = my_pipeline.transform(some_data)
```

```
[48]: model.predict(prepared_data)
```

```
[48]: array([20.717, 27.434, 12.339, 21.093, 22.174])
```

```
[49]: list(some_labels)
```

```
[49]: [20.9, 26.6, 12.0, 20.4, 22.2]
```

1.6 Evaluating the model

```
[50]: from sklearn.metrics import mean_squared_error
housing_predictions = model.predict(housing_num_tr)
mse = mean_squared_error(housing_labels, housing_predictions)
rmse = np.sqrt(mse)
```

```
[51]: rmse
```

```
[51]: 1.1751382402353208
```

```
[52]: #due to its high mean sq error we will discard this model and we'll use
      ↪decision tree regressor
      # but decision tree regressor does overfitting that's why mse is 0
```

1.7 Using Better Evaluation Technique- Cross Valdiation

```
[53]: # 1 2 3 4 5 6 7 8 9 10
from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, housing_num_tr, housing_labels, scoring=
      ↪'neg_mean_squared_error', cv= 10)
rmse_scores = np.sqrt(-scores)
```

```
[54]: rmse_scores
```

```
[54]: array([3.21732103, 2.53358929, 5.07545782, 2.7368161 , 2.69396689,
        2.53952812, 2.75901019, 2.99910886, 1.99235371, 4.29029414])
```

```
[55]: def print_scores(scores):
      print('Scores: ', scores)
      print('Mean: ', scores.mean())
      print('Std Dev: ', scores.std())
```

```
[56]: print_scores(rmse_scores)
```

```
Scores: [3.21732103 2.53358929 5.07545782 2.7368161 2.69396689 2.53952812
        2.75901019 2.99910886 1.99235371 4.29029414]
```

Mean: 3.0837446166561895
Std Dev: 0.8726628328940665

1.8 Saving the model

```
[57]: from joblib import dump, load
      dump(model, 'Dragon.joblib')
```

```
[57]: ['Dragon.joblib']
```

1.9 Testing the model on test data

```
[61]: x_test = strat_test_set.drop('MEDV', axis=1)
      y_test = strat_test_set['MEDV'].copy()
      x_test_prepared = my_pipeline.transform(x_test)
      final_predictions = model.predict(x_test_prepared)
      final_mse = mean_squared_error(y_test, final_predictions)
      final_rmse = np.sqrt(final_mse)
      print(final_predictions, list(y_test))
```

```
[22.774 22.234 46.452 32.708 45.269 34.833 20.967 23.491 32.813 19.749
 19.401 30.646 22.038 33.504 20.447 21.997 12.351 21.29 28.141 19.577
 20.035 45.194 12.158 19.231 26.17 34.225 16.545 15.826 6.566 20.523
 23.558 23.228 18.153 15.227 20.715 18.828 22.893 17.532 45.359 17.284
 21.484 18.683 19.441 18.604 33.154 8.203 24.735 14.436 21.167 21.333
 45.969 23.799 14.994 21.599 19.767 46.899 33.329 20.024 35.048 10.564
 23.8 35.34 33.224 23.76 14.238 20.778 21.026 15.79 27.983 24.352
 23.375 32.209 19.282 31.91 11.048 20.14 42.456 19.75 19.804 14.03
 41.829 8.996 35.687 23.043 28.801 15.961 23.127 21.926 20.614 16.192
 26.234 9.877 32.011 12.73 25.99 20.553 33.358 13.687 21.429 21.372
 20.813] [24.6, 22.0, 44.8, 23.6, 48.8, 36.5, 19.7, 23.1, 34.6, 21.5, 23.1,
 15.0, 23.0, 34.9, 18.5, 10.4, 10.2, 18.9, 23.9, 19.3, 19.4, 48.3, 10.9, 19.6,
 27.5, 37.3, 16.1, 15.2, 10.5, 21.4, 23.2, 20.7, 21.7, 13.0, 22.3, 19.6, 21.2,
 18.1, 50.0, 23.7, 22.6, 20.5, 18.9, 19.5, 32.7, 8.8, 29.1, 19.0, 22.6, 21.2,
 50.0, 22.5, 17.8, 20.3, 20.4, 37.6, 35.4, 18.2, 33.3, 12.1, 23.1, 37.9, 36.1,
 23.7, 13.1, 23.8, 19.6, 13.1, 27.9, 27.0, 22.9, 31.7, 17.1, 30.3, 8.1, 19.6,
 44.0, 19.5, 18.5, 17.2, 35.2, 8.3, 34.7, 20.5, 23.7, 14.2, 22.8, 20.6, 19.6,
 15.2, 23.9, 6.3, 32.0, 13.4, 22.0, 19.9, 28.7, 19.1, 23.4, 11.9, 21.7]
```

```
[59]: final_rmse
```

```
[59]: 3.4284325619272673
```

```
[63]: prepared_data[0]
```

```
[63]: array([-0.44241248,  3.18716752, -1.12581552, -0.27288841, -1.42038605,
        -0.54568298, -1.7412613 ,  2.56284386, -0.99534776, -0.57387797,
        -0.99428207,  0.43852974, -0.49833679])
```

1.10 Using the model

```
[64]: from joblib import dump, load
import numpy as np
model = load('Dragon.joblib')
features = np.array([[-0.44241248, 34.18716752, -1.12581552, -0.27288841, -1.
↪42038605,
                    -99.54568298, -122.7412613 , 9.56284386, -0.99534776, -0.57387797,
                    -0.99428207, 0.43852974, -9.49833679]])
model.predict(features)
```

```
[64]: array([23.91])
```

```
[ ]:
```